

# **Wids-2025 Mid-Term Report**

**FlappyBird: where Flappy learns to fly**

**Prashant Meena (24B0949)**

**04/01/2026**



# CHAPTER 1: WEEK 1&2

**J**N WEEK-1, I practiced Python and its libraries, including NumPy, Pygame, and Matplotlib.

In week-2, I studied Markov chains and Markov Decision Processes (MDP), focusing on state design, reward design, discretization, and environment testing.

The assignment solutions were uploaded to the GitHub repository: <https://github.com/prashantdlp/FlappyBird-AI>

## MARKOV CHAINS AND MARKOV DECISION PROCESSES

Many sequential decision-making problems can be modeled as a Markov Decision Process (MDP) under the assumption of full state observability. An MDP is defined as a 5-tuple  $(S, A, P, R, \gamma)$ , where  $S$  is the set of states,  $A$  is the set of actions,  $P$  denotes the state transition probabilities (i.e., the consequences of actions),  $R$  is the reward function providing feedback to the agent, and  $\gamma \in [0, 1]$  is the discount factor.

The discount factor  $\gamma$  determines the importance of future rewards. When  $\gamma = 0$ , the agent considers only immediate rewards, leading to greedy behavior. As  $\gamma$  approaches 1, future rewards are weighted more heavily, encouraging long-term planning. In episodic tasks,  $\gamma = 1$  implies that future and immediate rewards are valued equally.

### A BIT FORMAL :)

At each time step  $t$ , the agent observes the current state  $s_t \in S$  and selects an action  $a_t \in A$  according to a policy  $\pi(a_t | s_t) = P(a_t = a | s_t = s)$ . The environment then transitions to a new state  $s_{t+1}$  sampled from the transition distribution  $P(\cdot | s_t, a_t)$  and emits a scalar reward  $r_t = R(s_t, a_t)$ .

The return is defined as the discounted cumulative reward:

$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}.$$

The value function of a policy  $\pi$  is given by:

$$V^\pi(s) = \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right]$$

and the corresponding action-value function is:

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a \right]$$

The value function satisfies the Bellman expectation equation:

$$V^\pi(s) = \sum_{a \in A} \pi(a | s) \sum_{s' \in S} P(s' | s, a) [R(s, a) + \gamma V^\pi(s')].$$

The optimal value function is defined as:

$$V^*(s) = \max_\pi V^\pi(s),$$

and satisfies the Bellman optimality equation:

$$V^*(s) = \max_{a \in A} \sum_{s' \in S} P(s' | s, a) [R(s, a) + \gamma V^*(s')].$$

The objective of reinforcement learning is to find an optimal policy  $\pi^*$  that maximizes the expected discounted return:

$$\pi^* = \arg \max_\pi \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right].$$

## MARKOV PROPERTY

A state signal that succeeds in retaining all relevant information is said to be Markov, or to have the Markov property .

## DISCRETIZATION

Tabular Q-learning requires a finite set of states, whereas most real-world environments are continuous. Discretization addresses this mismatch by mapping continuous variables into a finite number of bins. This introduces an inherent trade-off. Very fine discretization leads to a large Q-table, slow convergence, and sparse updates, while overly coarse discretization results in faster learning at the cost of precision and suboptimal decisions. The objective is to choose a resolution that is sufficient for effective decision-making without unnecessary complexity.

**Flappy Bird Insight.** In Flappy Bird, the agent does not require exact pixel-level information. Instead, decision-relevant abstractions are sufficient:

- Whether the bird is above or below the pipe gap
- The direction and magnitude of its vertical velocity
- Whether the next pipe is near or far

Thus, discretization should capture task-relevant structure rather than perfect environmental accuracy.

## **REWARD DESIGN**

---

Reward design plays a crucial role in shaping agent behavior. The reward function should align closely with the true objective of the task while remaining simple and stable. Poorly designed rewards can lead to unintended behaviors, such as reward hacking or learning shortcuts that do not solve the intended problem. Sparse rewards make learning difficult due to limited feedback, whereas overly dense or large rewards can dominate learning and cause instability. Effective reward design balances guidance and flexibility, encouraging progress toward the goal without over-constraining the policy.

## **ENVIRONMENT VALIDATION AND STRESS TESTING**

---

Before training an agent, it is essential to verify that the environment is well-defined and learnable. Random policies provide a simple yet effective sanity check, helping identify implementation bugs, impossible objectives, or poorly scaled rewards. If a random agent fails immediately, the task may be overly difficult; if it survives indefinitely, the task may be trivial. Unbounded or exploding rewards often indicate reward design errors.

***Environment Stress Testing.*** Reinforcement learning systems can be highly sensitive to small changes in environment parameters. Stress testing evaluates robustness by varying key factors such as gravity, reward scaling, and observation noise. Agents that perform well only under a narrowly tuned configuration are brittle and lack generalization. Analyzing sensitivity to such perturbations provides insight into which environmental variables dominate the learned behavior.